

Channel-Aware Backdoor Attacks Against Federated Infostealer Malware Classification Using Dynamic API-Call and Network Representations

Mochamad Asryl Aziz

BPS-Statistics Indonesia, Bantaeng Regency

Moh. Jabir Mubarok

Institut Teknologi Sepuluh Nopember, Surabaya

Eka Fitria

Informatics, Universitas Syiah Kuala, Banda Aceh

2026 8th International Workshop on Big Data and Information Security (IWBIS)
16-18th October 2026

Background and Motivation

Infostealer malware

Steals credentials, browser data and session cookies [1], [2].

Knowing the family helps threat intelligence, triage and incident response.

Dynamic analysis

API calls show what the malware does on the host; network artifacts show how it communicates.

Both can be encoded as images for a CNN [3], [4], [5].

Federated learning

Clients train one classifier together without moving raw data.

A malicious client can inject a backdoor through its updates [6], [7].

Research Problem and Objectives

Research problem

Most FL backdoor work targets general image or model-level poisoning, not the channels of a malware representation [7].

RGB-stack places API-call, network and fused information in separate channels [5], so a trigger in one channel need not behave like a trigger in another.

Objectives

- Build API-call and network representations from Cuckoo Sandbox reports of five latest and most prevalent infostealer families.
- Select the representation and the CNN backbone.
- Evaluate channel-aware backdoors in FL with red, green, blue and full-RGB triggers, trigger controls and Multi-Krum.

The novelty is to treat the representation's channels as the attack surface: a trigger confined to the fusion channel alone matches one spanning all three.

Overall Research Workflow




Fig. 1. Overall research workflow of the proposed study.

From Sandbox Reports to Image Tiles and RGB-Stack




Fig. 2. Dataset creation pipeline from dynamic analysis reports to dynamic API-call and network representations.









	AgentTesla	FormBook	SalatStealer	StealC	Vidar
API-call sequences Representations	 16x16	 16x16	 16x16	 16x16	 16x16
Network Traffic Representations	 28x28	 28x28	 28x28	 28x28	 28x28
Opacity blend Fusion	 128x128	 128x128	 128x128	 128x128	 128x128
RGB-Stacks Fusion	 128x128	 128x128	 128x128	 128x128	 128x128

Fig. 3. Representative processed malware representations generated from dynamic API-call and network packet artifacts. Rows illustrate different representation types, while columns correspond to the five infostealer families.

Channel-Aware Backdoor in Federated Learning



Five clients, 70 images each, 14 per family: balanced and IID.

One malicious client puts a trigger on AgentTesla training images and relabels them as FormBook [7], [8].

Poison rate 20% of its 14 AgentTesla images: two images per round, resampled every round.

Four trigger settings: red/API, green/network, blue/fusion and full RGB.

Fig. 4. Federated learning architecture used in this study, consisting of one central server and five clients. One client is malicious in the backdoor attack setting.

Dataset and Experimental Environment

TABLE I. Dataset Summary and Stratified Split

Family	Total	Train	Val.	Test
AgentTesla	100	70	15	15
FormBook	100	70	15	15
SalatStealer	100	70	15	15
StealC	100	70	15	15
Vidar	100	70	15	15
Total	500	350	75	75

TABLE II. Experimental Environment

Component	Specification
Dynamic analysis	Cuckoo Sandbox on Ubuntu
Training environment	Windows, Python / Anaconda
Framework	PyTorch with CUDA 12.8
GPU	NVIDIA GeForce RTX 3080, 12 GB VRAM

Training, Defense and Evaluation

Training

Component	Specification
Batch size	16
Learning rates and weight decay	10^{-4}
FL server rounds	50
FL local rounds	2
Optimizer	AdamW

Model and evaluation

ResNet18 on RGB-stack, trained from scratch.

Final-round global model, with no best-validation selection.

Clean accuracy, macro-F1 and ASR over seeds 42, 123 and 2026.

Defense and trigger control

Multi-Krum is the primary defense [12].

Clipping, coordinate-wise median and trimmed mean are also screened [13].

Trigger control: the same triggers on the clean FL model at test time.

Backbone Selection Results

TABLE III. Backbone Selection

Representation	Backbone	Acc.	Macro-F1
Opacity blend	SmallCNN	0.5333	0.5298
Opacity blend	MobileNetV2	0.6933	0.6903
Opacity blend	ResNet18	0.7733	0.7730
RGB-stack	SmallCNN	0.5200	0.5144
RGB-stack	MobileNetV2	0.7200	0.7200
RGB-stack	ResNet18	0.7867	0.7884

Among the RGB-stack backbones, measured on seed 42, ResNet18 is the strongest: **78.67% accuracy and 78.84% macro-F1.**

It is used for all later experiments.

The opacity-blend rows are listed for reference only.

Channel-Aware Backdoor Sweep

TABLE IV. Channel-Aware Backdoor Sweep Using Seed 42

Trigger	Clean Acc.	Macro-F1	ASR
Red/API	0.8533	0.8524	0.3333
Green/network	0.8400	0.8419	0.3333
Blue/fusion	0.8400	0.8395	1.0000
Full RGB	0.8400	0.8401	1.0000

Blue/fusion and full RGB reach 100% ASR; red/API and green/network do not.

Blue/fusion does so without modifying all three channels.

Not a trigger artifact: blue/fusion and full RGB differ from their trigger controls at $p < 10^{-17}$ (Fisher exact); red and green do not ($p = 1, 5/15$ against 4/15).

Defense Screening

TABLE V. Single-Seed Defense Screening Against Strongest Attacks Using Seed 42

Trigger	Defense	Clean Acc.	Macro-F1	ASR
Blue/fusion	None	0.8400	0.8395	1.0000
Blue/fusion	Clipping	0.8533	0.8550	1.0000
Blue/fusion	Median	0.8133	0.8120	1.0000
Blue/fusion	Trimmed mean	0.8267	0.8244	1.0000
Blue/fusion	Multi-Krum	0.8267	0.8244	1.0000
Full RGB	None	0.8400	0.8401	1.0000
Full RGB	Clipping	0.8533	0.8550	1.0000
Full RGB	Median	0.8400	0.8388	1.0000
Full RGB	Trimmed mean	0.8133	0.8145	1.0000
Full RGB	Multi-Krum	0.7600	0.7542	0.4000

Clipping, coordinate-wise median and trimmed mean leave ASR at 1.0000.

Only Multi-Krum reduces it, and only under full RGB: 1.0000 to 0.4000.

Multi-Krum is taken to the multi-seed evaluation.

[12] P. Blanchard et al., NeurIPS, 2017; [13] D. Yin et al., ICML, 2018

Multi-Seed Results

TABLE VI. Main Multi-Seed Results Under IID Federated Setting, With Per-Seed Values

Scenario	Channel	Clean Acc.	Macro-F1	ASR / Target Rate	Per seed (42, 123, 2026)
Clean FL	—	0.8267 ± 0.0134	0.8255 ± 0.0104	—	—
Trigger control	Blue/fusion	0.8267 ± 0.0134	0.8255 ± 0.0104	0.1333 ± 0.1155	4/15, 1/15, 1/15
Trigger control	Full RGB	0.8267 ± 0.0134	0.8255 ± 0.0104	0.1556 ± 0.1018	4/15, 1/15, 2/15
Backdoor, FedAvg	Blue/fusion	0.8311 ± 0.0539	0.8320 ± 0.0539	1.0000 ± 0.0000	15/15, 15/15, 15/15
Backdoor, FedAvg	Full RGB	0.8267 ± 0.0611	0.8262 ± 0.0621	1.0000 ± 0.0000	15/15, 15/15, 15/15
Backdoor, Multi-Krum	Blue/fusion	0.8178 ± 0.0154	0.8184 ± 0.0135	0.5111 ± 0.4286†	15/15, 3/15, 5/15
Backdoor, Multi-Krum	Full RGB	0.7956 ± 0.0407	0.7931 ± 0.0427	0.6000 ± 0.3464†	6/15, 15/15, 6/15

† Across-seed ASR range ≥ 0.5: the outcome is bimodal, so read the per-seed counts (out of 15 source samples), not the mean.
The seed-42 clean baseline and trigger controls come from a 30-round, one-local-epoch run. No ASR result is affected.

Across three seeds, blue/fusion and full-RGB backdoors reach 100% ASR, with clean performance comparable to the clean FL baseline.

In seven of the eight trigger controls the target rate is identical with and without the trigger: the ASR is not an artifact of the pattern.

Defense Analysis: Bimodal Rather Than Partial




Fig. 7. Per-round clean macro-F1 and ASR across three random seeds for the blue/fusion and full-RGB backdoor settings under FedAvg and Multi-Krum. Shaded areas denote standard deviation across seeds.

Under Multi-Krum the mean is misleading. Per seed: 15/15, 3/15, 5/15 for blue/fusion and 6/15, 15/15, 6/15 for full RGB.

Bimodal: on some seeds the backdoor remains fully effective, on the others most of it is suppressed.

It tracks selection: the malicious client is retained in 8–50% of rounds, and that rate correlates with the final ASR ($r = 0.893$, $p = 0.017$, $n = 6$).

Multi-Krum is unreliable as a standalone defense.

The Blue/Fusion Channel

The ordering in Table IV

Mean value inside the trigger region of the clean images: 214.6 in R, 48.3 in G and 21.6 in B.

37.7% of red trigger-region pixels are already at least 254, where the trigger changes nothing.

Contrast explains red's failure but not green's, whose contrast is 207.

What it suggests, and what it does not

Blue is an edge map of the other two, so its effectiveness cannot come from information they lack.

It is almost empty, so the trigger is the only strong response there.

Consistent with the measurements, but not established.

Conclusion and Future Work

- **The representation's channels are an attack surface: a trigger confined to the fusion channel alone matches one spanning all three.**
- Blue/fusion and full-RGB triggers reached 100% ASR across three seeds, with clean performance close to the baseline.
- Multi-Krum reduced ASR bimodally rather than partially, leaving the backdoor fully effective (15/15) on one of the three seeds for each trigger.
- **Channel-aware backdoors are a serious threat to federated malware classifiers in this controlled IID setting, and motivate representation-aware defenses.**

Limitations and future work

Limitations: the IID-only partition, a single source–target pair, and three seeds, with the red and green results resting on seed 42 alone.

Future work: a non-IID partition, a second source–target pair and a contrast-matched trigger.

Thank you

Mochamad Asryl Aziz · Moh. Jabir Mubarak · Eka Fitria

mochamadasryl93@gmail.com · jabirmubarak@gmail.com · ekafitria2@mhs.usk.ac.id